

ALGOTHON 2026

Jump Trading Track — AlgoSoc Novice Category

1. Summary

Our team built a mathematically rigorous **trading platform** from the ground up. We deployed a **low-latency event-driven dispatcher** that routes every orderbook tick to one of five proprietary engines, each targeting a distinct mathematical edge across all eight synthetic London markets. Fair values are computed dynamically through **Rolling Multivariate Ridge Regressions**, **Rolling OLS Co-Integration**, and **Exponentially Weighted Moving Averages (EWMA)** meaning that the bot always adapts instead of statistically guessing.

Strategy	Markets Targeted	Mathematical Model
A — M2vsM1StatArb	TIDE_SPOT, TIDE_SWING	Rolling Ridge Regression Z-Score
B — ETFPackStatArb	LONETF, LONFLY	Ridge Regression + IV Spread
C — ETFBasketArb	LONETF (all 3 legs)	Identity Arbitrage (no model risk)
D — UpDipBuyer	TIDE_SPOT, WX_SUM, WX_SPOT	Asymmetric EWMA Dip Buyer
E — PairTradingArb	LHR_COUNT, LHR_INDEX	Rolling OLS Co-Integration

2. Advanced Mathematical Models

A. Rolling OLS Co-Integration Pairs Arbitrage

Target: Heathrow Airport Markets — M5: LHR_COUNT vs M6: LHR_INDEX

Markets 5 and 6 are both derived from exactly the same Heathrow PIHub flight dataset. M5 counts total movements; M6 measures net directional imbalance. This creates a **latent structural co-integration** that our engine exploits in real time.

The engine fits a **Rolling Ordinary Least Squares** regression over a 600-tick window to continuously estimate $M6 = \alpha + \beta(M5)$. The residual is normalised to a Z-score over a 350-tick lookback. When the Z-score breaches $\pm 2.2\sigma$, the engine executes an **asymmetric β -weighted delta hedge**: selling the rich leg and buying the cheap leg in a ratio proportional to the live OLS beta, keeping the combined position market-neutral.

Signal Condition	Action Taken	Risk Control
Z-score > $+2.2\sigma$	SELL M6 · BUY $\beta \times M5$	clip=12 per leg, max_pos=70
Z-score < -2.2σ	BUY M6 · SELL $\beta \times M5$	clip=12 per leg, max_pos=70
$ Z\text{-score} < 0.6\sigma$	FLATTEN both legs	Hard limit ± 100 (BotExchangeAdapter)

Engine Implementation

```

class PairTradingArb :
    """ Rolling OLS co-integration arb: M6 ≈ α + β · M5. """
    def step (self):
        a_mid = self.api.get_mid (self.pA)      # LHR_COUNT live mid
        b_mid = self.api.get_mid (self.pB)      # LHR_INDEX live mid

        # — Live OLS Fit —
        self._ols.add (a_mid , b_mid )
        _a , beta = self._ols.fit ()           # closed-form β

        # — Residual & Z-score —
        fair_b = self._ols.predict (a_mid )
        resid = b_mid - fair_b
        self._rz.add (resid )
        z = self._rz.z(resid )               # normalised residual

        # — β-weighted Hedge Sizing —
        hedge_a = int (np.clip (round (beta) ,
                               - self.hedge_clip , self.hedge_clip ))
        if z > self.z_enter and room_sell_b > 0:
            q_b = int (min (self.clip_B , room_sell_b ))
            self.api.place_order (self.pB , "SELL" , bid_b , q_b )
            q_a = int (min (self.clip_A , room_buy_a ,
                           abs (hedge_a ) * q_b ))
            self.api.place_order (self.pA , "BUY" , ask_a , q_a )

```

Figure A1 — `PairTradingArb.step()`: live OLS fit → residual → Z-score → β -weighted hedge execution

Visualising the Engine Mechanics




Figure A2 — Simulated LHR_COUNT × LHR_INDEX. Top: live M6 price vs OLS fair value with signal zones and execution markers. Middle: raw residual series. Bottom: normalised Z-score with $\pm 2.2\sigma$ entry bands and $\pm 0.6\sigma$ exit thresholds.

B. Asymmetric EWMA Dip Buyer


Target: $TIDE_SPOT \cdot WX_SUM \cdot WX_SPOT$

Tidal and weather products exhibit structural **upward drift** — the Thames approaches High Tide as the settlement window closes, and temperature \times humidity tends to peak mid-day. Symmetric mean-reversion is dangerous here: shorting artificial tops builds toxic short inventory into a rising market.

Our engine tracks an **Exponentially Weighted Moving Average** ($\alpha = 0.06$) and computes a rolling Z-score of the deviation. It strictly refuses to short, acting only on the long side when price dislocates below -1.6σ (*liquidity void / transient flush*) and taking profit when the Z-score recovers above -0.3σ .

Parameter	TIDE_SPOT	WX_SUM	WX_SPOT
EWMA α	0.06	0.06	0.06
Entry threshold	-1.6σ	-1.6σ	-1.6σ
Take-profit	-0.3σ	-0.3σ	-0.3σ
Max position	± 80	± 45	± 45
Clip per tick	12	8	8

Engine Implementation



```

class UpDipBuyer :
    """ EWMA mean-reversion - buys sharp dips, never shorts tops. """

    def step (self):
        mid = self.api.get_mid (self.product)
        bbo = self.api.get_best_bid_ask (self.product)

        # --- EWMA tracking & Rolling Volatility ---
        self._ema.update (mid)                      # α = 0.06
        dev = mid - self._ema.mu
        self._dev.add (dev)
        sd = self._dev.std ()

        # --- Z-score Normalisation ---
        z = dev / sd if not np.isnan (sd) else np.nan

        # --- Entry Logic - asymmetric BUY only ---
        if not np.isnan (z) and z < -self.z_enter:
            pos = self.api.get_position (self.product)
            room_buy = max (0, self.max_pos - pos)
            if room_buy > 0:
                q = int (min (self.clip, room_buy))
                self.api.place_order (
                    self.product, "BUY", bbo.best_ask, q)

        # --- Take-Profit Exit ---
        if not np.isnan (z) and z > -self.z_take:
            pos = self.api.get_position (self.product)
            if pos > 0:
                q = int (min (self.clip, pos))
                self.api.place_order (
                    self.product, "SELL", bbo.best_bid, q)

```

Figure B1 — UpDipBuyer.step(): EWMA tracking → Z-score normalisation → asymmetric BUY-only entry → take-profit exit

Visualising the Engine Mechanics




Figure B2 — Simulated TIDE_SPOT. Top: live price vs EWMA with buy-execution markers (\blacktriangle). Middle: signed deviation from EWMA. Bottom: Z-score with -1.6σ trigger line and -0.3σ take-profit line.

3. Engineering Safety & System Stability

Every order placed by any strategy passes first through our centralised **BotExchangeAdapter** — a bridge layer that enforces the ± 100 competition position limit as an absolute hard cap before touching the exchange. No strategy can breach this constraint regardless of signal strength or simultaneous strategy overlap.

- **Sub-Position Clipping:** each engine limits its injection per tick (clip=12 for arb legs, clip=8 for dip buyers), distributing liquidity gradually rather than wiping out entire orderbook levels.
- **Coldown Timers:** all five engines carry independent cooldown periods (0.15–0.2 s), preventing feedback loops where a partial fill immediately re-triggers the same signal.
- **WARMUP Guards:** no engine trades until its statistical estimators (OLS, Ridge, EWMA) have accumulated sufficient history. NaN propagation from early ticks is explicitly blocked with early-return guards.
- **MM Suppression:** passive market-making is disabled on LHR_COUNT and LHR_INDEX — both legs of the pair arb — eliminating self-churn between the MM module and PairTradingArb.
- **Continuous Integration:** 145 unit and integration tests in pytest validate every strategy, helper class, and dispatch path under mocked market conditions before any live deployment.

4. Live Data Pipeline Engineering

Fair-value estimation requires real-world London data ingested faster than the exchange tick rate. We built a parallel data pipeline fetching three independent external sources simultaneously on every refresh cycle, with per-source TTL caching to avoid hammering external APIs.

Data Source	API Endpoint	Markets Priced
Thames Water Level	UK Environment Agency (15-min)	M1 TIDE_SPOT, M2 TIDE_SWING
London Weather	Open-Meteo (51.5°N, 0.1°W)	M3 WX_SPOT, M4 WX_SUM

Heathrow Flights	Heathrow PIHub (official API)	M5 LHR_COUNT, M6 LHR_INDEX
Derived	M1 + M3 + M5 computed in-process	M7 LON ETF, M8 LON_FLY

5. Retrospective & Real-World Application

Why Our Architecture Underperformed:

- **The Spread-Crossing Tax:** Our Ridge and OLS engines correctly identify structural mispricings but must cross the bid/ask spread to act, instantly paying a 2-tick transaction cost that eliminates the mathematical edge in a low-liquidity simulator.
- **Synthetic Market Irrationality:** Co-integration models rely on markets eventually reverting to structural laws. In a novice tournament driven by thousands of random order injections, fundamental relationships can permanently decouple, turning mathematical edges into drawn-out losses.
- **Latency Equality:** In this simulator all bots operate on identical discrete event ticks, neutralising the HFT speed advantage these strategies depend on in live markets.

Why This Architecture Wins in Live Markets:

- **Making vs Taking Liquidity:** At institutional speed, our engines rest passive limit orders at the calculated fair value, capturing the spread rather than paying it — reversing the entire cost structure.
 - **Real-World Co-Integration:** Professional market makers enforce structural bounds across correlated instruments. When LHR_COUNT and LHR_INDEX misprice relative to their PIHub constraints, multi-billion-dollar firms ruthlessly arbitrage the gap, confirming the mean-reversion our OLS engine predicts.
 - **Adaptive Hedging:** Hardcoded algorithms blow up when volatility regimes shift. Our Rolling Ridge and OLS regressions recalculate β tick-by-tick, keeping the portfolio risk-neutral under macro shocks.
-